

The poweRlaw package

Colin S. Gillespie

Last updated: March 24, 2013

The powerLaw package provides code to fit discrete and continuous power-law distributions. The fitting procedure follows the method detailed in Clauset *et al.*¹. The scaling coefficient, α , is obtained by maximising the likelihood. The cut-off value, x_{\min} , is estimated by minimising the Kolmogorov-Smirnoff statistic.

Future versions of this package will allow other heavy tailed distributions to be fitted.

¹ A. Clauset, C.R. Shalizi, and M.E.J. Newman. Power-law distributions in empirical data. *SIAM review*, 51(4):661–703, 2009

1 Installation

The package is hosted on github.² The package can be installed using the devtools package:³

```
install.packages("devtools")
library(devtools)
install_github("powerLaw", "csgillespie", subdir = "pkg")
```

Once installed, the package can be loaded ready for use with the standard library command

```
library(powerLaw)
```

² The intention is eventually host this package on CRAN.

³ If use Windows, you need to install the Rtools package first.

2 Accessing documentation

I have tried to ensure that the package and all associated functions and datasets are properly documented with runnable examples. The command

```
help(package = "powerLaw")
```

will give a brief overview of the package and a complete list of all functions. The list of vignettes associated with the package can be obtained with

```
vignette(package = "powerLaw")
```

At the time of writing, *this* vignette is the only one available, and can be accessed from the R command line with

```
vignette("powerLaw", package = "powerLaw")
```

Help on functions can be obtained using the usual R mechanisms. For example, help on the function `rpldis` can be obtained with

```
?rpldis
```

and the associated example can be run with

```
example(rpldis)
```

A list of demos and data sets associated with the package can be obtained with

```
demo(package = "powerLaw")
data(package = "powerLaw")
```

For example, the Moby dick data set can be load using

```
data(moby)
```

After running this command, the vector `moby` will be accessible, and can be examined by typing

```
moby
```

at the R command prompt.

3 Example: Word frequency in Moby Dick

This example investigates the frequency of occurrence of unique words in the novel Moby Dick by Herman Melville⁴. The data can be downloaded from

<http://tuvalu.santafe.edu/~aaronc/powerlaws/data.htm>

or loaded directly

```
data(moby)
```

3.1 Fitting a discrete power-law

To fit a discrete power-law, we create a discrete power-law object, `displ`

```
m_m = displ$new(moby)
```

Initially the lower cut-off, x_{\min} is set to the smallest x value and the scaling parameter, α , is set to NULL

```
m_m$getXmin()
## [1] 1
```

```
m_m$getPars()
## NULL
```

The distribution object also has standard setters

The package also contains the data set `moby_sample`. This data set is 2000 randomly sampled values from the larger `moby` data set.

⁴ A. Clauset, C.R. Shalizi, and M.E.J. Newman. Power-law distributions in empirical data. *SIAM review*, 51(4):661–703, 2009; and M.E.J. Newman. Power laws, pareto distributions and zipf’s law. *Contemporary physics*, 46(5):323–351, 2005

```
m_m$setXmin(5)
m_m$setPars(2)
```

For a given x_{\min} value, we can estimate the corresponding α value using its maximum likelihood estimator (mle)

```
estimate_pars(m_m)
## [1] 1.921
```

To estimate the lower bound, we minimise the distance between the data and the fitted model CDF, that is

$$D(x) = \max_{x \geq x_{\min}} |S(x) - P(x)|$$

where $S(x)$ is the data CDF and $P(x)$ is the theoretical CDF. The value $D(x)$ is known as the Kolmogorov-Smirnov statistic. Our estimate of x_{\min} is then the value of x that minimises $D(x)$:

```
(est = estimate_xmin(m_m))
## $KS
## [1] 0.009229
##
## $xmin
## [1] 7
##
## $pars
## [1] 1.95
##
## attr("class")
## [1] "ks_est"
```

We can then set parameters of power-law distribution to the "optimal" values

```
m_m$setXmin(est)
```

All distribution objects have generic plot methods:⁵

```
## Plot the data (from xmin)
plot(m_m)
## Add in the fitted distribution
lines(m_m, col = 2)
```

When calling the plot and lines function, the data plotted is actually invisibly returned, i.e.

```
dd = plot(m_m)
head(dd, 3)
##   x     y
## 1 1 1.0000
## 2 2 0.5141
## 3 3 0.3505
```

Instead of using the mle, we could instead do a parameter scan:
`estimate_pars(m_m, pars=seq(2, 3, 0.1))`




Figure 1: Plot of the data CDF for the Moby Dick data set. This corresponds to figure 6.1(a) in Clausett, 2009. Plot of data CDF with line of best fit.

⁵ Generic lines and points functions are also available.

Algorithm 1: Uncertainty in x_{\min}

```

1: Set  $N$  equal to the number of values in the original data set
2: for  $i$  in  $1:B$ :
3:   Sample  $N$  values from the original data set
4:   Estimate  $x_{\min}$  and  $\alpha$  using the Kolmogorov-Smirnoff statistic.
5: end for

```

3.2 Uncertainty in x_{\min}

Clauset, *et al*, 2009 recommend a bootstrap procedure to get a handle on parameter uncertainty. Essentially, we sample with replacement from the data set and then re-infer the parameters.

To run the bootstrapping procedure, we use the `bootstrap` function

```
bs = bootstrap(m_m, no_of_sims = 1000, threads = 1)
```

this function runs in parallel, with the number of threads used determined by the `threads` argument. To detect the number of cores on your machine, you can run:

```
parallel::detectCores()
## [1] 4
```

The object return by `bootstrap` is a list with three elements:

- The original Kolmogorov-Smirnov statistic
- The results of the bootstrapping procedure
- The average time (in seconds) for a single bootstrap

The results of the bootstrap are best investigated with histograms

```
hist(bs$bootstraps[, 2], breaks = "fd")
hist(bs$bootstraps[, 3], breaks = "fd")
```

and a bivariate scatter plot

```
plot(bs$bootstraps[, 2], bs$bootstraps[, 3])
```

These commands produce figures 2 and 3.

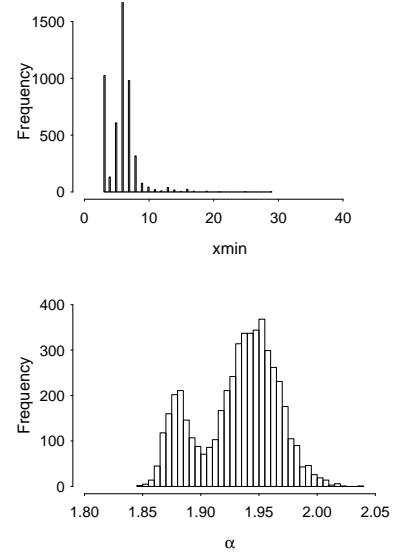


Figure 2: Characterising uncertainty in parameter values. (a) x_{\min} uncertainty (standard deviation 2) (b) α uncertainty (std dev. 0.03)




Figure 3: Characterising uncertainty in parameter values. Bivariate scatter plot of x_{\min} and α .

Algorithm 2: Do we have a power-law?

-
- ```

1: Calculate point estimates x_{\min} and the scaling parameter α .
2: Set n_1 equal to the number of values below x_{\min} .
3: Set $n_2 = n - n_1$.
4: for i in $1:B$:
5: Simulate n_1 values from a discrete uniform distribution: $U(1, x_{\min})$ and n_2 values from a discrete
 power-law (with parameter α).
6: Estimate x_{\min} and α using the Kolmogorov-Smirnoff statistic.
7: end for

```
- 

## 3.3 Do we have a power-law?

Clausett, *et al*, 2009 recommend a bootstrap procedure to determine whether the underlying distribution is a power-law the uncertainty. Essentially, we perform a hypothesis test by generating multiple data sets (with parameters  $x_{\min}$  and  $\alpha$ ) and then "re-inferring" the model parameters. The algorithm is detailed in Algorithm 2.<sup>6</sup>

When  $\alpha$  is close to one, this algorithm can be particularly time consuming to run, for two reasons:

1. When generating random numbers from the discrete power-law distribution, extreme values are highly possible, i.e. values greater than  $10^8$ . Hence, when generating the random numbers, all numbers larger than  $10^5$  are generated using a continuous approximation.
2. To calculate the Kolmogorov-Smirnov statistic, we need explore the state space. It is computationally infeasible to explore the entire state space when  $\max(x) \gg 10^5$ . So to this algorithm feasible, we explore two state space. The first,

$$x_{\min}, x_{\min} + 1, x_{\min} + 2, \dots, 10^5$$

and combine it with an additional  $10^5$  values from

$$10^5, \dots, \max(x)$$

The bootstrapping procedure, steps 4 – 7, can be run in parallel. To estimate the uncertainty with the moby data set, we use

```

##This may take a while
##Use the mle to estimate the parameters
bs_p = bootstrap_p(m_m, no_of_sims=1000, threads=1)

```

The object returned from the bootstrap procedure contains four elements

- A  $p$ -value - `bs$p`. See section 4.2 of the Clauset paper.
- The original goodness of fit statistic - `bs$gof`.
- The result of the bootstrap procedure - a data frame with three columns.
- The average time (in seconds) for a single bootstrap realisation.

The results of this procedure are shown in figures 4 and 5.

<sup>6</sup> Algorithm 2 can easily be extended for other distributions.

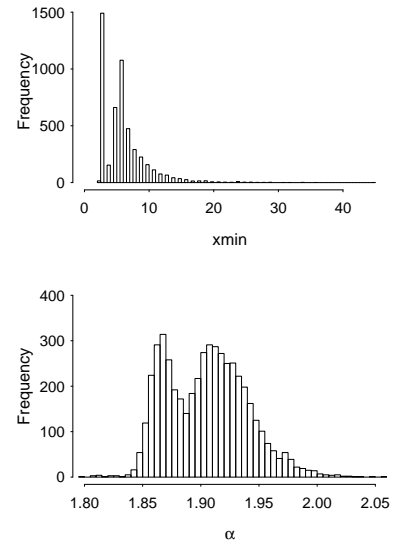


Figure 4: Histograms of the bootstrap results.




Figure 5: Bivariate scatter plot of the bootstrap results. The values of  $x_{\min}$  and  $\alpha$  are obviously strongly correlated.

## 4 Distribution objects

During the Moby Dick example, we created a `displ` object

```
m_m = displ$new(moby)
```

The object `m` has class `displ` and inherits the general distribution class. A list of available distributions is given in table 1.

| Distribution       | Object name        | # Parameters |
|--------------------|--------------------|--------------|
| Discrete Power-law | <code>displ</code> | 1            |
| CTN Power-law      | <code>conpl</code> | 1            |

Table 1: Available distributions in the power-law package. These objects are all reference classes.

All distribution objects list in table 1 are reference classes. The key point, is that unlike S4 classes, reference classes have a mutable state. Each distribution object has four fields:

- `datatype`: This will be set to *discrete* or *continuous*.
- `dat`: a copy of the data.
- `xmin`: the lower cut-off  $x_{\min}$ .
- `pars`: a vector of parameter values.
- `internal`: a list of values use in different numerical procedures. This will differ between distribution objects.

By using the mutable state, we have efficient caching of data structures that can be reused. For example, the mle of discrete power-laws uses the statistic:

$$\sum_{i=x_{\min}}^n \log(x_i)$$

This value is calculated once for all values of  $x_{\min}$ , then iterated over when estimating  $x_{\min}$ .

All distribution objects have a number of methods available. A list of methods is given in table 2. See the associated help files for further details.

| Method Name                | Description                                                                     |
|----------------------------|---------------------------------------------------------------------------------|
| <code>dist_cdf</code>      | Cumulative density/mass function (CDF)                                          |
| <code>dist_pdf</code>      | Probability density/mass function (pdf)                                         |
| <code>dist_rand</code>     | Random number generator                                                         |
| <code>dist_data_cdf</code> | Data CDF                                                                        |
| <code>dist_ll</code>       | Log-likelihood                                                                  |
| <code>estimate_xmin</code> | Point estimates of the cut-off point and parameter values                       |
| <code>estimate_pars</code> | Point estimates of the parameters (conditional on the current $x_{\min}$ value) |
| <code>bootstrap</code>     | Bootstrap procedure (uncertainty in $x_{\min}$ )                                |
| <code>bootstrap_p</code>   | Bootstrap procedure to test whether we have a power-law                         |

Table 2: A list of functions for distribution functions. These objects do not change the object states. However, they may not be thread safe.

See `?setRefClass` for further details on references classes.

## 5 Loading data

Typically, data is stored in a csv or text file. To use this data, we load it in the usual way<sup>7</sup>

```
blackouts = read.table("blackouts.txt")
```

Distribution objects take vectors as inputs, so

```
m_bl = conpl$new(blackouts$V1)
```

## 6 Comparison with the *plfit* script

### 6.1 The discrete case

Other implementations of estimating the lower bound can be found at

<http://tuvalu.santafe.edu/~aaronc/powerlaws/>

In particular, the script for estimating  $x_{\min}$  can be loaded using

```
source("http://tuvalu.santafe.edu/~aaronc/powerlaws/plfit.r")
```

The results are directly comparable to the *powerLaw* package. For example, if we look consider the Moby Dick data set again:

```
plfit(moby)
$xmin
[1] 7
##
$alpha
[1] 1.95
##
$D
[1] 0.009289
```

Notice that the results are slightly different. This is because the *plfit* by default does a parameter scan over the range

1.50, 1.51, 1.52, ..., 2.49, 2.50

To exactly replicate the results, we could use

```
estimate_xmin(m_m, pars = seq(1.5, 2.5, 0.01))
```

### 6.2 The continuous case

The *plfit* script also fits continuous power-laws. Again the results are comparable.

For example, suppose we have one thousand random numbers from a continuous power-law distributinos with parameters  $\alpha = 2.5$  and  $x_{\min} = 10.0$

<sup>7</sup> The blackouts data set can be obtained from Clauset's website: <http://goo.gl/BsqnP>.




Figure 6: CDF plot of the blackout dataset with line of best fit. Since the minimum value of  $x$  is large, we fit a continuous power-law as this is more efficient.

```
r = rplcon(1000, 10, 2.5)
```

The `plfit` automatically detects if the data is continuous

```
plfit(r)
$xmin
[1] 13.41
##
$alpha
[1] 2.595
##
$D
[1] 0.02456
```

Fitting with the `powerLaw` package gives the same values

```
m_r = conpl$new(r)
(est = estimate_xmin(m_r))
$KS
[1] 0.02456
##
$xmin
[1] 13.41
##
$pars
[1] 2.595
##
attr("class")
[1] "ks_est"
```

Of course, using the `powerLaw` package we can easily plot the data

```
m_r$setXmin(est)
plot(m_r)
lines(m_r, col = 2)
```

to get figure 7.




Figure 7: CDF plot of one thousand random numbers generated from a power-law with parameters  $\alpha = 2.5$  and  $x_{\min} = 10$ . The line of best fit is also shown.



## *References*

- [1] A. Clauset, C.R. Shalizi, and M.E.J. Newman. Power-law distributions in empirical data. *SIAM review*, 51(4):661–703, 2009.
- [2] M.E.J. Newman. Power laws, pareto distributions and zipf’s law. *Contemporary physics*, 46(5):323–351, 2005.

*Package and R version*

| Package  | Version |
|----------|---------|
| parallel | 2.15.3  |
| powerLaw | 0.16.1  |
| VGAM     | 0.8-4   |

Table 3: A list of packages and versions used.

```

version
##
platform x86_64-pc-linux-gnu
arch x86_64
os linux-gnu
system x86_64, linux-gnu
status
major 2
minor 15.3
year 2013
month 03
day 01
svn rev 62090
language R
version.string R version 2.15.3 (2013-03-01)
nickname Security Blanket

```